Correlation-Aware Belief Aggregation: Deterministic Spatial-Temporal Discounting for Byzantine-Tolerant Sensor Fusion over Merkle-DAGs

Fabrizio Salmi

Independent Researcher

March 2026 — Technical Report / Preprint

Abstract. The AIMP Epistemic Layer (L3) aggregates distributed evidence using integer log-odds arithmetic and two-pass Markovian trust propagation, achieving 98–142× speedup over Subjective Logic and Dempster-Shafer with bit-identical results across architectures. However, L3 v0.2.0 treats every evidence source as statistically independent — a Naive Bayes assumption that produces pathological hyper-confidence when physically correlated sensors (e.g., co-located IoT devices) or semantically correlated agents (e.g., LLMs fine-tuned on the same dataset) report concordant observations.

We present *Grid-Cell Correlation Discounting*, an extension that partitions evidence by discrete correlation coordinates and applies geometric decay within each correlated cluster. The strongest source retains full weight; each subsequent source receives $\text{discount}_{\text{bps}}^{k}/10000^{k}$ of its original contribution. We solve the CRDT associativity challenge — geometric decay is inherently non-associative across partial merges — by extending the existing grid-aligned epoch reduction to bucket by the triple (temporal epoch, semantic fingerprint, correlation cell). This guarantees that discounting is always computed atomically on the complete stabilized set, eliminating the associativity requirement.

The implementation adds 120 lines to the existing 3,000-line Rust codebase. All 50 existing tests pass unchanged; 14 new tests cover discounting correctness, commutativity, backward compatibility, and cell-isolated reduction. Claims with no correlation cell behave identically to v0.2.0 (zero regression). The implementation is open-source (Rust, MIT license).

1. Introduction

Decentralized sensor fusion faces a fundamental tension between *coverage* (more sensors improve accuracy) and *redundancy* (correlated sensors inflate confidence). In Bayesian aggregation, independent evidence combines via addition in log-odds space [1]: n sensors each reporting 70% confidence yield a combined posterior that grows monotonically with n, approaching certainty. When the sensors are independent, this is correct. When they are correlated — sharing a physical environment, a common data source, or a training distribution — the posterior is inflated beyond what the actual information content justifies.

This is not a theoretical concern. In Decentralized Physical Infrastructure Networks (DePIN), hundreds of sensors may occupy the same building, the same street corner, or the same drone swarm. In multi-agent LLM systems, models fine-tuned on overlapping corpora produce systematically correlated inferences. The AIMP Epistemic Layer v0.2.0 [2] explicitly documents this as *Limitation #1*: "Naive Bayes independence assumption — correlated sources produce hyper-confidence."

We address this limitation with three contributions:

1. **Grid-Cell Correlation Discounting.** We introduce a discrete correlation coordinate (`CorrelationCell(u64)`) that applications assign to claims based on spatial proximity (geohash), semantic similarity (model family ID), or temporal co-occurrence. Within each cell, evidence is ranked

by magnitude and geometrically discounted. All arithmetic is integer-only (basis points), preserving determinism and ZK-readiness.

2. **Atomic Cell Reduction.** We extend the existing grid-aligned epoch reduction to bucket claims by the triple (temporal epoch, semantic fingerprint, correlation cell). This architectural choice eliminates the associativity requirement for the discount function, because all nodes compute the discount on the identical complete set within each bucket — partial merges across cell boundaries never occur.
3. **Security Analysis.** We analyze three attack vectors specific to correlation metadata (Byzantine correlation claims, correlation stuffing, hyper-discounting) and show that the existing reputation system provides defense-in-depth.

1.1. Relationship to Prior Work

This paper builds directly on the AIMP protocol stack:

- **L1/L2 (v0.1.0)**: Merkle-DAG CRDT with Ed25519 signing, Noise Protocol transport, and BFT quorum consensus [3].
- **L3 (v0.2.0)**: Epistemic Layer with integer log-odds, two-pass Markovian trust propagation, Sybil-resistant reputation, and materialized compaction [2].

The present work modifies L3′s aggregation pipeline. All other components (trust propagation, cycle detection, contradiction damping, holographic routing) are unchanged.

Correlation-aware aggregation has been studied in the context of copulas [4], which model joint distributions of correlated random variables using continuous CDFs. While mathematically elegant, copula-based approaches require floating-point arithmetic and produce architecture-dependent results — incompatible with the BFT determinism requirement. Dempster-Shafer Theory [5] handles evidential correlation through mass function combination, but at 142× the computational cost of L3′s integer approach (v0.2.0 benchmark). Subjective Logic [6] addresses correlation via *dependent opinions*, but its floating-point operators produce non-deterministic results across architectures.

Our approach trades continuous expressiveness for discrete determinism: correlation is modeled as cell membership (binary: same cell or not) rather than pairwise distance (continuous). This is a deliberate engineering choice that preserves all CRDT and BFT invariants.

2. Problem Statement

2.1. The Naive Bayes Assumption in L3 v0.2.0

L3 v0.2.0 aggregates evidence via pure log-odds addition:

$$\text{aggregate}(e_1, ..., e_n) = \sum_{i=1}^{n} e_i$$

where e_i is the reputation-weighted log-odds of claim i. The only deduplication is by `evidence_source` hash (BLAKE3 of original data), which protects against *network-level* echo chambers (the same datum rebroadcast by different nodes) but not against *physical* or *semantic* correlation.

Example. Consider 100 temperature sensors on the same rooftop, each reporting 40°C with 70% confidence ($e_i = 847$ milli-log-odds). Each has a unique `evidence_source` (different hardware). L3 v0.2.0 aggregates:

$$\text{aggregate} = 100 \times 847 = 84700$$

This corresponds to a posterior probability indistinguishable from 100% — yet a single independent measurement from a different location would contribute only 847. The 100 co-located sensors provide essentially one observation's worth of independent information, not 100.

2.2. Design Constraints

Any solution must preserve the invariants established in v0.2.0:

Invariant	Source	Implication
Determinism	L3 core	No floats. All math in i32/i64.
CRDT safety	Reducer	Commutative, associative, idempotent.
BFT tolerance	Trust model	Correlation metadata can be Byzantine.
AP convergence	EEC	Must not break eventual convergence.
L2 independence	Architecture	L3 never blocks L2 merges.
ZK-readiness	Design goal	Bounded integer arithmetic only.

3. Grid-Cell Correlation Discounting

3.1. Correlation Cell

We define a *correlation cell* as a discrete coordinate that partitions the observation space into regions of expected dependence:

```
#[derive(Clone, Copy, PartialEq, Eq, Hash, PartialOrd, Ord)]
pub struct CorrelationCell(pub u64);
```

Cell assignment is application-defined, following the same philosophy as edge generation in v0.2.0 (the protocol specifies the mechanism; the application specifies the policy):

- **IoT/DePIN**: Geohash [7] truncated to N characters, cast to `u64`. A 6-character geohash covers approximately 1.2 km × 0.6 km — large enough to capture roof-level correlation while fine enough to distinguish city blocks.
- **Multi-agent LLM**: Model family identifier (e.g., Llama=1, Mistral=2). Models from the same family share training data and systematic biases.
- **Temporal**: `tick / temporal_grid_size`. Sensors sampling within the same temporal window observe the same physical event.
- **Composite**: `BLAKE3(spatial_cell || model_family || temporal_bucket)`.

The `Claim` struct gains an optional field:

```
pub correlation_cell: Option<CorrelationCell>,
```

Claims with `None` behave identically to v0.2.0 (backward compatibility).

3.2. Geometric Discount Factor

Within a correlated cell, claims are first *ranked* by absolute evidence strength (descending), with the cryptographic claim ID as deterministic tie-breaker. This ranking is critical: the *strongest* signal in the cluster always survives at full weight, ensuring that the most informative observation dominates. Weaker, redundant signals receive progressively less weight.

The weight assigned to rank-i claim (0-indexed) is:

$$w_i = e_i \times \frac{\text{discount}_{\text{bps}}^i}{10000^i}$$

where e_i is the reputation-weighted log-odds of the claim at rank i, and $\text{discount}_{\text{bps}} \in [0, 10000]$ is a protocol parameter (default: 3000 = 30%). The rank-0 claim receives 100% of its original weight; rank-1 receives 30%; rank-2 receives 9%; and so on.

3.2.1. Integer Implementation (No Floats)

The exponentiation $\text{discount}_{\text{bps}}^i / 10000^i$ cannot be computed directly without floating-point. We implement it as an iterative integer multiplication with truncation at each step:

```
pub fn discount_factor(rank: u32, discount_bps: u16) -> u64 {
    let mut factor: u64 = 10000; // 100% in basis points
```

```
    for _ in 0..rank.min(MAX_DISCOUNT_DEPTH) {
        factor = factor * (discount_bps as u64) / 10000;
    }
    factor
}
```

The key property of this loop is that the division `/ 10000` occurs *inside* each iteration, not after the final product. This prevents `u64` overflow: the maximum intermediate value is $10000 \times 10000 = 10^8$, well within `u64` range (1.8×10^{19}).

Each iteration truncates toward zero (Rust integer division semantics). After 10 iterations with $\text{discount}_{\text{bps}} = 3000$, the factor reaches 0 and stays there. `MAX_DISCOUNT_DEPTH = 30` caps the loop to prevent adversarial claims with artificially high rank from causing disproportionate computation.

Exact decay table for $\text{discount}_{\text{bps}} = 3000$:

Rank	Factor (bps)	Effective %
0	10000	100.0%
1	3000	30.0%
2	900	9.0%
3	270	2.7%
4	81	0.81%
5	24	0.24%
6	7	0.07%
7	2	0.02%
8	0	0.00%

After rank 8, all subsequent claims contribute exactly zero. The geometric series converges: $\sum_{i=0}^{\infty} (0.3)^i = 1/(1-0.3) \approx 1.43$. Thus, N correlated sensors contribute at most $1.43 \times$ the evidence of a single sensor, regardless of N.

3.2.2. Deterministic Sorting

Within each correlation cell, claims are sorted by:

1. **Absolute magnitude of log-odds** (descending): The strongest evidence gets rank 0 (full weight). This ensures the most *informative* signal dominates, not the most *recent* or *first-received*.
2. **Cryptographic claim ID** (ascending): When two claims have identical absolute confidence, the tie is broken by BLAKE3 hash. This is a deterministic function of the claim content, independent of message arrival order, gossip topology, or node identity. Two nodes with the same claims *always* produce the same ranking.

```
group.sort_by(|(lo_a, id_a), (lo_b, id_b)| {
    lo_b.0.abs().cmp(&lo_a.0.abs())  // strongest first
        .then_with(|| id_a.cmp(id_b)) // BLAKE3 tiebreak
});
```

This sorting is the mechanism that guarantees *commutativity*: the aggregate of claims $\{A, B, C\}$ is identical regardless of the order in which they arrive at the node.

3.3. Correlation-Aware Aggregation Algorithm

The complete aggregation procedure:

```
Input:  evidence[] with (logodds, cell, claim_id)
Output: discounted aggregate LogOdds

1. Group by CorrelationCell.
   - None → uncorrelated singleton (full weight, v0.2.0 behavior).
```

```
   - Some(cell) → correlated group.

2. For each correlated group G with |G| = k claims:
   a. Sort by (|logodds| desc, claim_id asc).
   b. For rank i ∈ [0, k):
      w_i = logodds_i × discount_factor(i, discount_bps) / 10000
   c. group_total = Σ_i w_i  (using i64::saturating_add)

3. For uncorrelated claims (cell = None):
   uncorrelated_total = Σ logodds_j  (using i64::saturating_add)

4. final = (uncorrelated_total + Σ group_totals)
           .clamp(SAFE_MIN, SAFE_MAX)
```

Overflow safety: All intermediate additions use `i64::saturating_add`. Under adversarial claim flooding (billions of `LogOdds::MAX` claims), the accumulator saturates at `i64::MAX` (9.2×10^{18}), is then clamped to `SAFE_MAX` (10^9), and the node continues operating. No panic path exists in the aggregation loop.

4. Atomic Cell Reduction: Solving the CRDT Associativity Problem

4.1. The Problem: Geometric Decay is Not Associative

The CRDT SemanticReducer requires three algebraic properties: *commutativity*, *associativity*, and *idempotency*. Geometric discounting satisfies commutativity (via deterministic sorting) and idempotency (via dedup-by-ID), but *not* associativity across partial merges.

4.1.1. Numerical Demonstration

Let A, B, C be three claims in the same `CorrelationCell`, all with confidence 1000 milli-log-odds and identical reputation. Let $\text{discount}_{\text{bps}} = 5000$ (50%).

Case 1: Batch reduction (all three at once).

The Reducer sees $\{A, B, C\}$, sorts by absolute magnitude (all equal, so tie-broken by ID: $A < B < C$), and applies geometric decay:

$$w_A = 1000 \times \frac{10000}{10000} = 1000 \qquad \text{(rank 0: 100\%)}$$

$$w_B = 1000 \times \frac{5000}{10000} = 500 \qquad \text{(rank 1: 50\%)}$$

$$w_C = 1000 \times \frac{2500}{10000} = 250 \qquad \text{(rank 2: 25\%)}$$

$$\text{batch}_{\text{total}} = 1000 + 500 + 250 = 1750$$

Case 2: Partial merge (asynchronous CRDT split-brain).

Node 1 is offline and sees only $\{A, B\}$. It reduces them into Summary_{AB}:

$$w_A = 1000, \quad w_B = 500 \quad \rightarrow \quad \text{Summary}_{AB}.\text{confidence} = 1500$$

The network reconnects. Node 2 must merge Summary_{AB} (confidence 1500) with claim C (confidence 1000). The Reducer now sees $\{\text{Summary}_{AB}, C\}$ and applies discounting:

$$w_{\text{Summary}} = 1500 \times \frac{10000}{10000} = 1500 \qquad \text{(rank 0)}$$

$$w_C = 1000 \times \frac{5000}{10000} = 500 \qquad \text{(rank 1)}$$

$$\text{partial}_{\text{total}} = 1500 + 500 = 2000$$

Violation:

$$1750 \neq 2000$$

The partial merge inflates the aggregate by 14.3%. Under repeated asynchronous compactions with different partition patterns, the final belief depends on *when* the network partitioned — a fatal violation of CRDT convergence.

4.2. The Solution: Architectural Atomicity via Triple Bucketing

We solve this *at the protocol level* rather than the mathematical level. The insight is that if partial merges *cannot occur* within a correlation cell, then the non-associative discount function is always applied to the complete set — and associativity becomes irrelevant.

4.2.1. Grid-Aligned Epoch Reduction (v0.2.0, Recap)

In v0.2.0, claims are bucketed by temporal epoch for deterministic reduction:

$$\text{bucket}_{\text{key}} = \left\lfloor \frac{\text{tick}}{\text{grid}_{\text{size}}} \right\rfloor$$

Two nodes that independently compact the same epoch produce byte-identical Summaries (same BLAKE3 hash), which L2 deduplicates. This guarantees that compaction is never applied to a *partial* temporal window.

4.2.2. Extended Bucketing (v0.3.0)

We extend the bucket key to include the correlation cell:

$$\text{bucket}_{\text{key}} = \left(\left\lfloor \frac{\text{tick}}{\text{grid}_{\text{size}}} \right\rfloor, \quad \text{correlation}_{\text{cell}} \right)$$

```
let mut buckets: BTreeMap<(u64, Option<u64>), Vec<Claim>> =
    BTreeMap::new();
for claim in claims {
    let epoch = claim.tick / grid_size;
    let cell_key = claim.correlation_cell.map(|c| c.0);
    buckets.entry((epoch, cell_key))
           .or_default()
           .push(claim.clone());
}
```

This guarantees four properties:

1. **Cell isolation.** All claims in a bucket share the same epoch *and* the same correlation cell. Claims from different cells are *never* mixed in the same reduction.
2. **Byte-identical Summaries.** Two nodes compacting the same bucket produce the same BLAKE3 hash, because the input set is identical (same epoch + same cell) and the discount function is deterministic (same sorting, same geometric factors).
3. **Atomic discount computation.** The discount is always computed on the *complete* set of claims in the bucket. There is no scenario where Node 1 discounts $\{A, B\}$ and Node 2 later merges the result with C — because C belongs to the same bucket and would have been included in the reduction on both nodes.
4. **Associativity is not required.** Since partial merges across cell boundaries are structurally impossible, the SemanticReducer's associativity requirement is satisfied *vacuously*: the reduction function is never called on a partial subset of a cell's claims within an epoch.

4.2.3. Summary Inheritance

The produced Summary inherits the `correlation_cell` of its input claims and includes it in the BLAKE3 hash:

```
if let Some(cell) = summary_cell {
    hasher.update(&cell.0.to_le_bytes());
}
// ...
Some(Claim {
    correlation_cell: summary_cell,
    // ...
})
```

Summaries from different cells are *cryptographically distinct*, even if they share the same semantic fingerprint and temporal epoch. This prevents cross-cell confusion during subsequent reductions.

4.2.4. Proof of CRDT Safety

Theorem (Atomic Discount Convergence). If two nodes N_1 and N_2 hold the same set of claims C after L2 convergence, they produce identical BeliefState vectors.

Proof. L2 convergence guarantees $C_1 = C_2 = C$. The grid-aligned reduction partitions C by $\left(\left\lfloor \frac{\text{tick}}{\text{grid}_{\text{size}}} \right\rfloor, \text{cell}\right)$. Within each bucket B_j:

(a) The set of claims is identical on both nodes: $B_j^1 = B_j^2$.

(b) Claims are sorted by (|confidence| desc, ID asc) — both are deterministic functions of claim content, independent of arrival order.

(c) The discount function $f(\text{sorted claims}, \text{discount}_{\text{bps}})$ is a pure function of its inputs (integer arithmetic, no state).

(d) Therefore $\text{Summary}_j^1 = \text{Summary}_j^2$ (byte-identical).

(e) L2 CRDT deduplicates identical Summaries natively (same hash).

(f) Trust propagation (unchanged from v0.2.0) operates on the converged claim set, producing identical BeliefState vectors. □

5. Integration with L3

5.1. Insertion Point

Correlation discounting is applied at two points in the L3 pipeline:

1. **`reduce_with_reputation()`**: The Reducer calls `LogOdds::aggregate_correlated()` instead of `LogOdds::aggregate()` when producing Summaries during epoch compaction.
2. **Base trust computation**: Future integration point for applying cell-aware discounting before trust propagation (Pass 1).

Trust propagation itself (Markovian flow, contradiction damping, temporal decay) is *unchanged*. The two-pass algorithm operates on the discounted base trust values, preserving the EEC proof from v0.2.0.

5.2. Backward Compatibility

Claims with `correlation_cell: None` follow the uncorrelated path in `aggregate_correlated()`, which performs pure `saturating_add` — identical to `LogOdds::aggregate()` from v0.2.0. The test `test_none_cell_backward_compat` verifies this: three claims at 2000 log-odds each produce 6000, exactly as in v0.2.0.

6. Security Analysis

6.1. Attack 1: Byzantine Correlation Evasion

Attack. A malicious node sets `correlation_cell = None` on all claims to bypass discounting, presenting correlated evidence as independent.

Defense (Topological Integrity). The `CorrelationCell` is included in the `Claim`'s BLAKE3 hash. Setting `None` does not merely change a metadata annotation — it creates a *topologically different claim* with a different cryptographic identity. This has two consequences:

(a) *Hash divergence*: If an honest node produces the same observation with a valid cell, the two claims have different IDs. The `evidence_source` dedup does not merge them, so the attacker does not gain additional weight — they produce a *separate* claim that must stand on its own reputation.

(b) *Contradiction detection*: Honest anchor nodes observing the same physical environment can emit `Contradicts` edges against claims that omit or falsify their correlation cell. The two-pass trust propagation (Pass 2) applies the contradiction penalty, scaled by the anchor's reputation. After repeated Byzantine behavior, the attacker's reputation is slashed to zero via `ReputationEvent::EquivocationDetected`, and all future claims contribute zero evidence (Reputation $0 \times$ any confidence $= 0$).

In permissioned networks, operators can enforce `correlation_cell` as a required field at the L2 validation layer. In open networks, absence of correlation attestation defaults to maximum correlation (conservative discounting).

6.2. Attack 2: Correlation Stuffing (Spatial Sybil)

Attack. An attacker creates 10,000 Sybil nodes, each assigned to a *unique* `CorrelationCell`, to bypass cell-based discounting. Each cell contains exactly one claim, so no geometric decay is applied.

Defense (Reputation Economics). This attack collides with the existing Sybil defense from v0.2.0. The attack fails in three layers:

1. **Layer 1: Zero-start reputation.** All 10,000 Sybil nodes start at Reputation 0. The `reduce_with_reputation()` function filters out claims from zero-reputation authors *before* grouping by cell:

   ```
   if rep.bps() == 0 {
       continue; // Zero-reputation authors excluded entirely
   }
   ```

   None of the Sybil claims survive to the discounting stage.

2. **Layer 2: Delegation spending.** To grant non-zero reputation to Sybils, the attacker must delegate from an established anchor. Delegation costs half of the granted reputation. Granting 2000 bps to each Sybil costs 1000 bps per delegation. After 10 delegations, the anchor's remaining reputation is $10000 - 10 \times 1000 = 0$. Total delegatable reputation converges to $2 \times$ original (geometric series bound from v0.2.0).

3. **Layer 3: Reputation weighting.** Even if a Sybil receives delegation, its evidence is scaled by $\text{rep} \times \frac{\text{confidence}}{10000}$. A 1-bps Sybil contributing $\text{VERY}_{\text{HIGH}} = 6907$ produces $6907 \times \frac{1}{10000} \approx 0.69$ milli-log-odds. Negligible.

Correlation discounting is a *second layer* of defense on top of reputation economics. The attacker must first defeat reputation gating (exponentially expensive) before cell assignment even matters.

6.3. Attack 3: Hyper-Discounting via Grid Collision

Attack. Honest sensors near a cell boundary are assigned to the same cell as a large cluster, causing legitimate independent observations to be unfairly discounted.

Defense (Tunable Resolution). Grid resolution is a protocol parameter, not a constant. Application operators choose the granularity:

- **Coarse grids** (geohash-4, 40 km): Strong discounting, may over-correlate independent sensors.
- **Fine grids** (geohash-7, 150 m): Weak discounting, rarely over-correlates, but may under-discount physically adjacent sensors.
- **Hierarchical grids** (future work): Multiple resolution levels with weighted combination.

The sensitivity analysis in the Evaluation section quantifies the tradeoff.

6.4. Attack 4: Integer Overflow via Claim Flooding

Attack. An attacker floods billions of claims with `LogOdds::MAX` ($\approx 2.14 \times 10^9$) values to cause `i64` overflow panic in the aggregation loop.

Defense (Saturating Arithmetic). All additions in `aggregate_correlated()` use `i64::saturating_add`:

```
total = total.saturating_add(lo.0 as i64);  // uncorrelated
total = total.saturating_add(discounted);     // correlated
```

Under adversarial flooding, the accumulator saturates docilely at `i64::MAX` (9.2×10^{18}), is then clamped to `SAFE_MAX` (10^9), and the node continues operating. There is no `panic!` path in the aggregation loop. This defense was added during multi-AI adversarial review (Round 1, overflow analysis).

7. Formal Verification

7.1. TLA+ Extension

We extend the `AimpBeliefConvergence.tla` specification with:

- A `CorrelationCell` model variable (domain: $\{1, 2, 3\}$).
- Modified `Aggregate` action with cell-aware grouping.
- New safety invariant **CorrelationBound**: the aggregate of N correlated claims must not exceed the aggregate of 1 full-weight claim plus $(N - 1)$ claims at `DISCOUNT_BPS` geometric decay.

7.2. Exhaustive Bounded Verification

The v0.2.0 exhaustive verifier (199,902 configurations, up to $N = 6$ nodes) is extended to cover correlation scenarios:

- Mixed correlated/uncorrelated claim sets.
- Cell boundary transitions during epoch compaction.
- Summary inheritance of correlation cells.

8. Evaluation

8.1. Test Suite

Test	Verifies
`test_uncorrelated_aggregate_unchanged`	None cells = v0.2.0 behavior
`test_same_cell_discounted`	5 correlated < 5 independent
`test_different_cells_independent`	Distinct cells = full weight
`test_mixed_correlated_and_independent`	Correct cell-by-cell math
`test_discount_commutativity`	Order-independent
`test_discount_idempotency`	Deterministic on duplicates
`test_epoch_aligned_splits_by_cell`	Separate Summaries per cell
`test_epoch_aligned_cell_isolation`	3+3 claims → 2 Summaries
`test_reduce_with_reputation_applies_discount`	Discount < naive sum
`test_none_cell_backward_compat`	Zero regression vs v0.2.0
`test_discount_factor_geometric_decay`	Exact bps values
`test_discount_factor_full_weight`	10000 bps = no discount
`test_discount_factor_zero_kills_all`	0 bps = rank-0 only
`test_geometric_decay_bounds`	Factor bounded at depth 30

All 64 tests pass (50 existing + 14 new).

8.2. Sensitivity Analysis: Discount Factor × Cluster Size

We compute the aggregate milli-log-odds for N correlated sensors, each with confidence 847 (70%), at varying `discount_bps` values:

N	Naive (v0.2.0)	10% discount	30% discount	50% discount	70% discount
1	847	847	847	847	847
5	4,235	939	1,205	1,638	2,347
10	8,470	939	1,207	1,687	2,739
50	42,350	939	1,207	1,687	2,809
100	84,700	939	1,207	1,687	2,809

Key observations:

1. The naive approach grows linearly with N: 100 sensors produce 100 × the evidence. This is the hyper-confidence pathology.
2. With 30% discount (default), the aggregate converges to 1,207 after approximately 8 claims (integer truncation yields slightly less than the theoretical $\frac{847}{1-0.3} \approx 1210$). Additional sensors beyond rank 8 contribute exactly zero.
3. With 70% discount (weak), the convergence is slower but still bounded at 2,809 (theoretical: $\frac{847}{1-0.7} \approx 2823$; integer truncation compounds across 30+ ranks).
4. With 10% discount (aggressive), convergence is immediate at 939 (theoretical: $\frac{847}{1-0.1} \approx 941$).
5. For $N \geq 10$, the discount effectively saturates. The aggregate is insensitive to cluster size beyond the convergence point.

8.3. Scenario: Mixed Correlated and Independent Sources

Consider a network with 3 independent cells of 10 sensors each, plus 5 fully independent sensors (no cell):

$$\text{total} = \underbrace{3 \times 1207}_{\text{3 cells, 30\% discount}} + \underbrace{5 \times 847}_{\text{5 independent}} = 3621 + 4235 = 7856$$

Compare with the naive approach:

$$\text{naive} = (3 \times 10 + 5) \times 847 = 35 \times 847 = 29645$$

The corrected aggregate (7856) is 3.8 × lower than naive (29645), reflecting the actual information content of 8 effective independent sources rather than 35 spurious ones.

8.4. Computational Overhead

The discount computation adds three operations to the reduction pipeline:

1. **BTreeMap grouping**: $O(n \log n)$ for n claims. For 1,000 claims with 10 cells, this adds approximately 50 ns (BTreeMap insertion on sorted keys).
2. **Intra-cell sorting**: $O(k \log k)$ for k claims per cell. For $k = 100$, approximately 200 ns (Rust `sort_by` on `i32` + `[u8;32]`).
3. **Discount factor computation**: $O(\min(r, 30))$ integer multiplications per claim. For rank 10, approximately 3 ns (10 `u64` multiply-divide pairs).

Total overhead estimate for 1,000 claims across 10 cells: approximately 300 ns. Compare with the v0.2.0 hot-path profile:

Component	v0.2.0 Time	v0.3.0 Overhead
Base trust computation	7.8%	+0 (unchanged)

Cycle detection (DFS)	31.6%	+0 (unchanged)
Trust propagation (2-pass)	59.0%	+0 (unchanged)
Classification	1.6%	+0 (unchanged)
Reduction (aggregate)	(included above)	+300 ns

For a 1,000-claim belief pipeline (140 µs total in v0.2.0), the correlation overhead is $300\frac{\text{ns}}{140}\mu\text{s} \approx 0.2\%$. The performance promise of v0.2.0 is preserved.

9. Limitations and Future Work

1. **Grid boundary effects.** Claims near cell boundaries may be incorrectly grouped. Hierarchical cells (coarse + fine grid) could mitigate this at the cost of additional complexity.
2. **Continuous distance metrics.** The discrete cell model cannot capture gradual correlation decay with distance. A future L3.1 could introduce pairwise distance-based discounting, though this requires $O(n^2)$ computation and careful treatment of floating-point determinism (likely via fixed-point distance).
3. **Cross-modal correlation.** A temperature sensor and an LLM analyzing satellite imagery of the same fire are correlated but occupy different cell types. Cross-modal correlation requires application-level fusion policies not addressed here.
4. **Contradiction attenuation within cells.** The discount function attenuates *all* weak signals within a cell, including contradictions. A correlated contradiction (e.g., the same faulty sensor cluster reporting "fire=false") is correctly weakened. However, an independent contradiction that happens to share a geohash cell with a strong positive cluster is also weakened — a false negative risk. The mitigation is that contradictions are primarily handled by the two-pass trust propagation (Pass 2), which operates *after* discounting and uses the global knowledge graph, not cell-local data.
5. **Dynamic discount parameter.** The `discount_bps` is a protocol constant. Adaptive discounting based on observed cell statistics (e.g., higher discount for denser cells) is a natural extension.
6. **Associativity at saturation boundaries.** The `SAFE_MAX` clamping behavior documented in v0.2.0 (Section 14) applies to the discounted aggregate as well. At extreme values, sequential and batch aggregation may produce different results. This is a pre-existing limitation, not introduced by correlation discounting.

10. Conclusion

We presented Grid-Cell Correlation Discounting for the AIMP Epistemic Layer, addressing the most critical limitation of v0.2.0 — the Naive Bayes independence assumption. The solution is deliberately conservative: discrete cells, integer arithmetic, and architectural atomicity via grid-aligned bucketing. It adds 120 lines of code, preserves all existing invariants, and introduces zero regression for uncorrelated claims.

The key architectural insight is that the CRDT associativity challenge can be solved at the *protocol level* (atomic bucketing) rather than the *mathematical level* (finding an associative discount function). This avoids the complexity of copula-based approaches [4] while providing equivalent protection in the discrete grid domain.

The sensitivity analysis demonstrates that with the default 30% discount, N correlated sensors converge to $\approx 1.42\times$ the evidence of a single sensor (1,207 vs 847 milli-log-odds) after rank 8 — eliminating the $N\times$ amplification of the naive approach while preserving full weight for independent observations from distinct cells.

Combined with the existing defenses (reputation gating, evidence source dedup, Markovian flow normalization), correlation discounting makes L3 the first BFT-deterministic belief aggregation framework that

addresses both network-level Sybil attacks and physical-level sensor correlation — entirely in integer arithmetic, entirely without coordination, and ready for zero-knowledge proof circuits.

Bibliography

[1] E. T. Jaynes, *Probability Theory: The Logic of Science*. Cambridge University Press, 2003.

[2] F. Salmi, "Byzantine-Tolerant Belief Aggregation over Merkle-DAGs: A Deterministic Epistemic Layer for Decentralized AI," 2026.

[3] F. Salmi, "AIMP: AI Mesh Protocol—Design and Evaluation of a Serverless Merkle-CRDT Protocol for Edge Agent Synchronization," 2026. [Online]. Available: https://www.researchgate.net/publication/403127328

[4] R. B. Nelsen, *An Introduction to Copulas*, 2nd ed. Springer, 2006.

[5] G. Shafer, *A Mathematical Theory of Evidence*. Princeton University Press, 1976.

[6] A. Jøsang, *Subjective Logic: A Formalism for Reasoning Under Uncertainty*. Springer, 2016.

[7] G. Niemeyer, "Geohash," 2008. [Online]. Available: https://en.wikipedia.org/wiki/Geohash